

Diamond Price Prediction Model

March 9, 2022

```
[1]: import sklearn

import os
import pandas as pd
import numpy as np

import shap
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import OneHotEncoder, LabelEncoder
import warnings
warnings.filterwarnings("ignore")

%matplotlib inline
```

```
[ ]:
```

```
[2]: #Load the Data
diamonds = pd.read_csv("C:/Users/Owner/Downloads/wholesale_diamonds.csv")
```

```
[3]: #a short preview of our dataset setructure
diamonds.head()
```

```
[3]:
```

	index	carat	cut	color	clarity	depth	table	cost (dollars)	\
0	0	0.23	Ideal	E	SI2	61.5	55.0	326	
1	1	0.23	Good	E	VS1	56.9	65.0	327	
2	2	0.29	Premium	I	VS2	62.4	58.0	334	
3	3	0.31	Good	J	SI2	63.3	58.0	335	
4	4	0.24	Very Good	J	VVS2	62.8	57.0	336	

	length (mm)	width (mm)	height (mm)	year
0	3.95	3.98	2.43	2010
1	4.05	4.07	2.31	2010
2	4.20	4.23	2.63	2010
3	4.34	4.35	2.75	2010
4	3.94	3.96	2.48	2010

```
[4]: #a little more preview of the dataset
diamonds.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 407280 entries, 0 to 407279
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   index                  407280 non-null int64
1   carat                  405232 non-null float64
2   cut                    407280 non-null object
3   color                  407280 non-null object
4   clarity                407280 non-null object
5   depth                  407280 non-null float64
6   table                  407280 non-null float64
7   cost (dollars)         407280 non-null int64
8   length (mm)            407280 non-null float64
9   width (mm)             407280 non-null float64
10  height (mm)            407280 non-null float64
11  year                   407280 non-null int64
dtypes: float64(6), int64(3), object(3)
memory usage: 37.3+ MB
```

Noticed that the index column is unnecessary and need to to drop

```
[5]: #drop the index column
diamonds = diamonds.drop(columns={"index"})
```

```
[6]: diamonds.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 407280 entries, 0 to 407279
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   carat                  405232 non-null float64
1   cut                    407280 non-null object
2   color                  407280 non-null object
3   clarity                407280 non-null object
4   depth                  407280 non-null float64
5   table                  407280 non-null float64
6   cost (dollars)         407280 non-null int64
7   length (mm)            407280 non-null float64
8   width (mm)             407280 non-null float64
9   height (mm)            407280 non-null float64
10  year                   407280 non-null int64
dtypes: float64(6), int64(2), object(3)
```

memory usage: 34.2+ MB

```
[7]: #Renaming the columns name
#columns:Cost,Length,Width and height and summarize statistics
diamonds.rename(columns ={'cost (dollars)': 'price', 'length (mm)': 'x', 'width (mm)': 'y', 'height (mm)': 'z'}, inplace = True)
diamonds.describe()
```

```
[7]:
```

	carat	depth	table	price \
count	405232.000000	407280.000000	407280.000000	407280.000000
mean	0.797742	61.747793	57.457113	4372.968506
std	0.474774	1.434209	2.239837	4503.620949
min	0.200000	43.000000	43.000000	-998.000000
25%	0.400000	61.000000	56.000000	1043.000000
50%	0.700000	61.800000	57.000000	2655.000000
75%	1.040000	62.500000	59.000000	5960.000000
max	4.130000	79.000000	95.000000	26930.000000

	x	y	z	year
count	407280.000000	407280.000000	407280.000000	407280.000000
mean	5.730165	5.732369	3.538519	2015.500000
std	1.122960	1.114266	0.712168	3.452057
min	0.000000	0.000000	0.000000	2010.000000
25%	4.710000	4.720000	2.910000	2012.750000
50%	5.690000	5.710000	3.520000	2015.500000
75%	6.530000	6.530000	4.030000	2018.250000
max	10.140000	10.100000	31.800000	2021.000000

```
[8]: # Price is int64, best if all numeric attributes have the same datatype,
      ↪especially as float64
diamonds["price"] = diamonds["price"].astype(float)
```

```
[9]: #preview the data set again
diamonds.head()
```

```
[9]:
```

	carat	cut	color	clarity	depth	table	price	x	y	z	year
0	0.23	Ideal	E	SI2	61.5	55.0	326.0	3.95	3.98	2.43	2010
1	0.23	Good	E	VS1	56.9	65.0	327.0	4.05	4.07	2.31	2010
2	0.29	Premium	I	VS2	62.4	58.0	334.0	4.20	4.23	2.63	2010
3	0.31	Good	J	SI2	63.3	58.0	335.0	4.34	4.35	2.75	2010
4	0.24	Very Good	J	VVS2	62.8	57.0	336.0	3.94	3.96	2.48	2010

0.0.1 Cleaning the Data

```
[10]: #check if we have a null value
diamonds.isnull().sum()
```

```
[10]: carat      2048
      cut        0
      color      0
      clarity    0
      depth      0
      table      0
      price      0
      x          0
      y          0
      z          0
      year       0
      dtype: int64
```

Noticed there are 2048 null values in Carat. They need to be dropped

```
[11]: #Drop the null values
      diamonds = diamonds.dropna(axis = 0, how='any')
```

```
[12]: #check if the null values are removed or not
      diamonds.isnull().sum()
```

```
[12]: carat      0
      cut        0
      color      0
      clarity    0
      depth      0
      table      0
      price      0
      x          0
      y          0
      z          0
      year       0
      dtype: int64
```

```
[13]: diamonds.isnull().values.any()
```

```
[13]: False
```

WE Noticed there are negative values in price. They need to be dropped to 0

```
[14]: #drop negative values to 0
      diamonds.drop(diamonds[diamonds['price'] <= 0].index, inplace = True)
      diamonds.describe()
```

```
[14]:
```

	carat	depth	table	price \
count	403192.000000	403192.000000	403192.000000	403192.000000
mean	0.797700	61.747838	57.457640	4397.571519
std	0.474791	1.434470	2.240196	4501.692426
min	0.200000	43.000000	43.000000	304.000000

25%	0.400000	61.000000	56.000000	1053.000000
50%	0.700000	61.800000	57.000000	2677.500000
75%	1.040000	62.500000	59.000000	5983.000000
max	4.130000	79.000000	95.000000	26930.000000

	x	y	z	year
count	403192.000000	403192.000000	403192.000000	403192.000000
mean	5.730006	5.732210	3.53842	2015.500243
std	1.123061	1.114364	0.71237	3.452026
min	0.000000	0.000000	0.00000	2010.000000
25%	4.710000	4.720000	2.91000	2013.000000
50%	5.690000	5.710000	3.52000	2016.000000
75%	6.530000	6.530000	4.03000	2019.000000
max	10.140000	10.100000	31.80000	2021.000000

```
[15]: diamonds.shape
```

```
[15]: (403192, 11)
```

```
[16]: diamonds = diamonds.drop(columns={"year"})
```

```
[17]: diamonds.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 403192 entries, 0 to 407279
Data columns (total 10 columns):
#   Column      Non-Null Count  Dtype
---  -
0   carat       403192 non-null  float64
1   cut         403192 non-null  object
2   color       403192 non-null  object
3   clarity     403192 non-null  object
4   depth       403192 non-null  float64
5   table       403192 non-null  float64
6   price       403192 non-null  float64
7   x           403192 non-null  float64
8   y           403192 non-null  float64
9   z           403192 non-null  float64
dtypes: float64(7), object(3)
memory usage: 33.8+ MB
```

```
[18]: diamonds.head()
```

```
[18]:   carat      cut color clarity depth table price     x     y     z
0   0.23    Ideal     E    SI2   61.5   55.0  326.0  3.95  3.98  2.43
1   0.23     Good     E    VS1   56.9   65.0  327.0  4.05  4.07  2.31
2   0.29  Premium     I    VS2   62.4   58.0  334.0  4.20  4.23  2.63
3   0.31     Good     J    SI2   63.3   58.0  335.0  4.34  4.35  2.75
```

```
4    0.24  Very Good      J    VVS2    62.8    57.0    336.0    3.94    3.96    2.48
```

```
[19]: diamonds.shape
```

```
[19]: (403192, 10)
```

0.0.2 Exploring and Visualization the data

It's easier to work a dataset when all its attributes are numerical. The cut, color and clarity attributes are non-numeric (They are objects). We still have to convert them to be numerical. Let's find out what categories exist for each of them.

```
[20]: # The diamond cut categories
diamonds["cut"].value_counts()
```

```
[20]: Ideal          161196
Premium         102668
Very Good       90813
Good            36701
Fair           11814
Name: cut, dtype: int64
```

```
[21]: #The diamond color categories
diamonds["color"].value_counts()
```

```
[21]: G           85245
E           72607
F           70154
H           63045
D           50803
I           40276
J           21062
Name: color, dtype: int64
```

```
[22]: # The diamond clarity categories
diamonds["clarity"].value_counts()
```

```
[22]: SI1          97900
VS2          91040
SI2          69326
VS1          60653
VVS2         37631
VVS1         28089
IF           13083
I1           5470
Name: clarity, dtype: int64
```

Let's take a preview of the summary of the numerical attributes and then an histogram on the dataset.

```
[23]: # Summary of each numerical attribute
diamonds.describe()
```


```
[23]:
```

	carat	depth	table	price \
count	403192.000000	403192.000000	403192.000000	403192.000000
mean	0.797700	61.747838	57.457640	4397.571519
std	0.474791	1.434470	2.240196	4501.692426
min	0.200000	43.000000	43.000000	304.000000
25%	0.400000	61.000000	56.000000	1053.000000
50%	0.700000	61.800000	57.000000	2677.500000
75%	1.040000	62.500000	59.000000	5983.000000
max	4.130000	79.000000	95.000000	26930.000000

	x	y	z
count	403192.000000	403192.000000	403192.000000
mean	5.730006	5.732210	3.53842
std	1.123061	1.114364	0.71237
min	0.000000	0.000000	0.000000
25%	4.710000	4.720000	2.91000
50%	5.690000	5.710000	3.52000
75%	6.530000	6.530000	4.03000
max	10.140000	10.100000	31.80000

```
[24]: diamonds.to_csv("C:/Users/Owner/Desktop/diamond_set/new_diamonds.csv")
```

```
[25]: diamonds.hist(bins = 50, figsize = (20, 15))
plt.show()
```




0.0.3 Understanding the correlation between variables

We have learned that to avoid a sampling bias, it is a good practice to perform stratified sampling. Stratified sampling is a technique that divides the dataset into homogeneous subgroups called strata. To use this, we need a right attribute of the dataset to predict the price. To get the right attribute, we need to pick an attribute most correlated to price. To select the correct attribute, we use the standard correlation coefficient

```
[26]: # Create a correlation matrix between every pair of attributes
corr_matrix = diamonds.corr()

# Plot the correlation with seaborn
cmap = sns.diverging_palette(70,20,s=50, l=40, n=6,as_cmap=True)
corrmat= diamonds.corr()
f, ax = plt.subplots(figsize=(12,12))
sns.heatmap(corrmat,cmap=cmap,annot=True, )
```

```
[26]: <AxesSubplot:>
```

Carat has the strongest correlation with price with the number of 0.92 X,y,z also have strong correlations with price depth and table have the weakest correlation with price

```
[27]: diamonds.corr()
```

```
[27]:
```

	carat	depth	table	price	x	y	z
carat	1.000000	0.027497	0.188253	0.915346	0.976104	0.975667	0.946206
depth	0.027497	1.000000	-0.294142	-0.013748	-0.025652	-0.028533	0.094620
table	0.188253	-0.294142	1.000000	0.134365	0.200761	0.194485	0.153990
price	0.915346	-0.013748	0.134365	1.000000	0.880170	0.882150	0.849546
x	0.976104	-0.025652	0.200761	0.880170	1.000000	0.998137	0.962398
y	0.975667	-0.028533	0.194485	0.882150	0.998137	1.000000	0.961687

```
z      0.946206  0.094620  0.153990  0.849546  0.962398  0.961687  1.000000
```

```
[28]: diamonds.describe()
```

```
[28]:
```


	carat	depth	table	price \
count	403192.000000	403192.000000	403192.000000	403192.000000
mean	0.797700	61.747838	57.457640	4397.571519
std	0.474791	1.434470	2.240196	4501.692426
min	0.200000	43.000000	43.000000	304.000000
25%	0.400000	61.000000	56.000000	1053.000000
50%	0.700000	61.800000	57.000000	2677.500000
75%	1.040000	62.500000	59.000000	5983.000000
max	4.130000	79.000000	95.000000	26930.000000

	x	y	z
count	403192.000000	403192.000000	403192.000000
mean	5.730006	5.732210	3.53842
std	1.123061	1.114364	0.71237
min	0.000000	0.000000	0.00000
25%	4.710000	4.720000	2.91000
50%	5.690000	5.710000	3.52000
75%	6.530000	6.530000	4.03000
max	10.140000	10.100000	31.80000

We see that carat correlates best with price. Its score is pretty high! Now we use this for our Stratified Sampling.

Let's take a closer look at the carat's histogram.

```
[29]: diamonds["carat"].hist(bins = 50)
plt.show()
```



```
[30]: diamonds.describe()
```

```
[30]:
```

	carat	depth	table	price \
count	403192.000000	403192.000000	403192.000000	403192.000000
mean	0.797700	61.747838	57.457640	4397.571519
std	0.474791	1.434470	2.240196	4501.692426
min	0.200000	43.000000	43.000000	304.000000
25%	0.400000	61.000000	56.000000	1053.000000
50%	0.700000	61.800000	57.000000	2677.500000
75%	1.040000	62.500000	59.000000	5983.000000
max	4.130000	79.000000	95.000000	26930.000000

	x	y	z
count	403192.000000	403192.000000	403192.000000
mean	5.730006	5.732210	3.53842
std	1.123061	1.114364	0.71237
min	0.000000	0.000000	0.000000
25%	4.710000	4.720000	2.91000
50%	5.690000	5.710000	3.52000
75%	6.530000	6.530000	4.03000
max	10.140000	10.100000	31.80000


```
[31]: # Divide the diamond carats by 0.4 to limit the number of carat categories
# Round up to have discrete categories
diamonds["carat_cat"] = np.ceil(diamonds["carat"] / 0.35)
```

```
# Merge categories > 5 in 5
diamonds["carat_cat"].where(diamonds["carat_cat"] < 5, 5.0, inplace = True)
```

```
[32]: # Check the distribution of the diamonds in the categories
diamonds["carat_cat"].value_counts()
```

```
[32]: 2.0    120609
      3.0    100673
      1.0     82624
      4.0     50871
      5.0     48415
      Name: carat_cat, dtype: int64
```

```
[33]: diamonds["carat_cat"].hist()
plt.show()
```



```
[34]: # Import the sklearn module
from sklearn.model_selection import StratifiedShuffleSplit

# Run the split. Creates on split and shares 20% of the dataset for the test set
split = StratifiedShuffleSplit(n_splits = 1, test_size = 0.2, random_state = 42)

# Separate the stratified train set and the test set
for x_train, x_test in split.split(diamonds, diamonds["carat_cat"]):
    strat_train_set = diamonds.iloc[x_train]
    strat_test_set = diamonds.iloc[x_test]
```

```
[35]: for set in (strat_train_set, strat_test_set):
        set.drop(["carat_cat"], axis = 1, inplace = True)
```

```
[36]: # Redefined diamonds dataset
diamonds = strat_train_set.copy()
diamonds.head()
```

```
[36]:
```

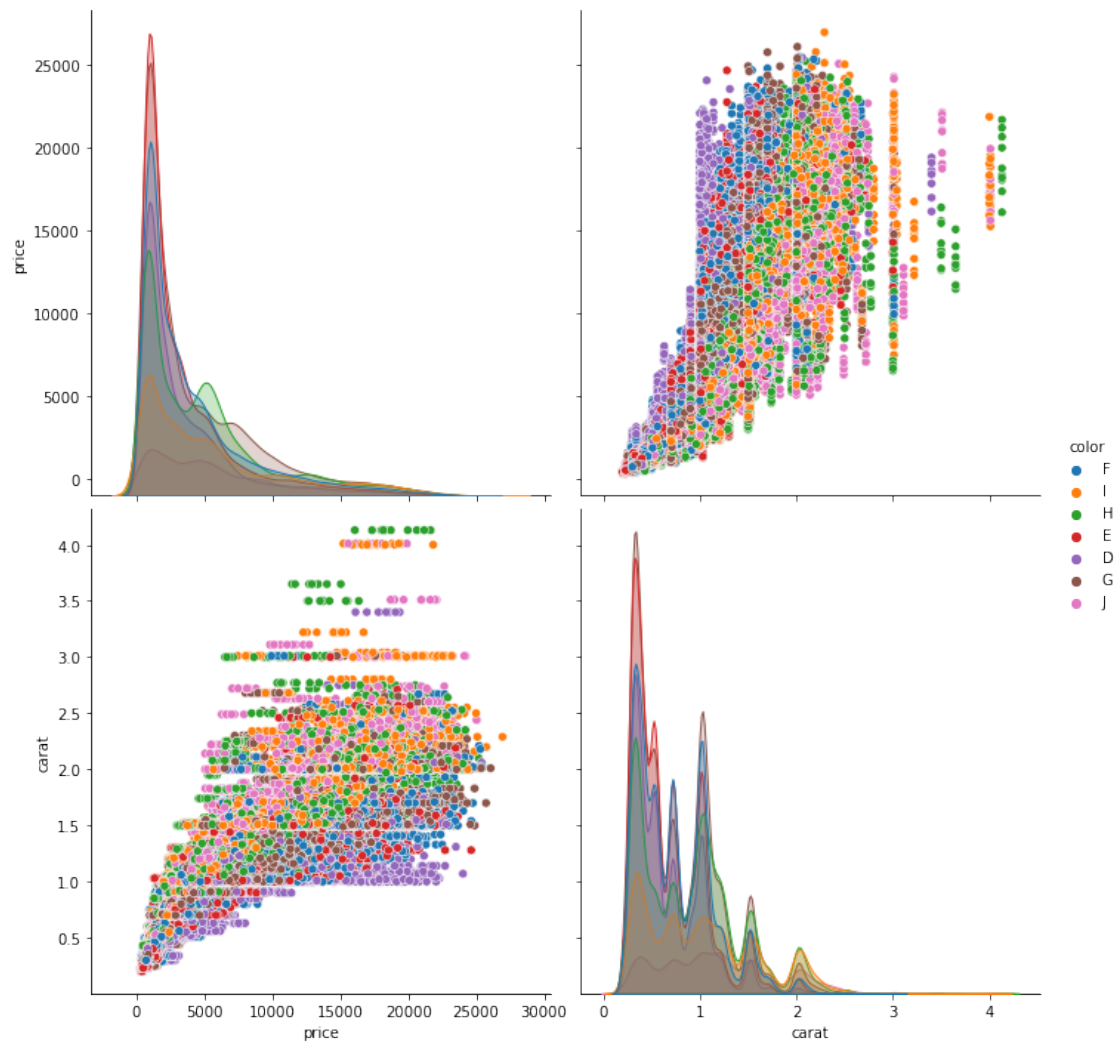
	carat	cut	color	clarity	depth	table	price	x	y	z
79064	1.09	Ideal	F	VS2	61.2	56.0	7616.0	6.68	6.65	4.08
80834	1.50	Very Good	I	VS2	60.1	60.0	9339.0	7.38	7.35	4.43
106815	1.19	Premium	H	SI2	61.9	58.0	4525.0	6.84	6.77	4.21
269312	0.75	Premium	I	SI1	62.0	60.0	2744.0	5.80	5.75	3.57
8251	1.21	Premium	H	VS2	60.8	62.0	5461.0	6.82	6.78	4.13



We can now perform a Stratified Sampling based on the carat categories : I will use Scikit-Learn's StratifiedShuffleSplit class.

0.1 Data Visualization

0.1.1 Let's play around with visualization of the data set

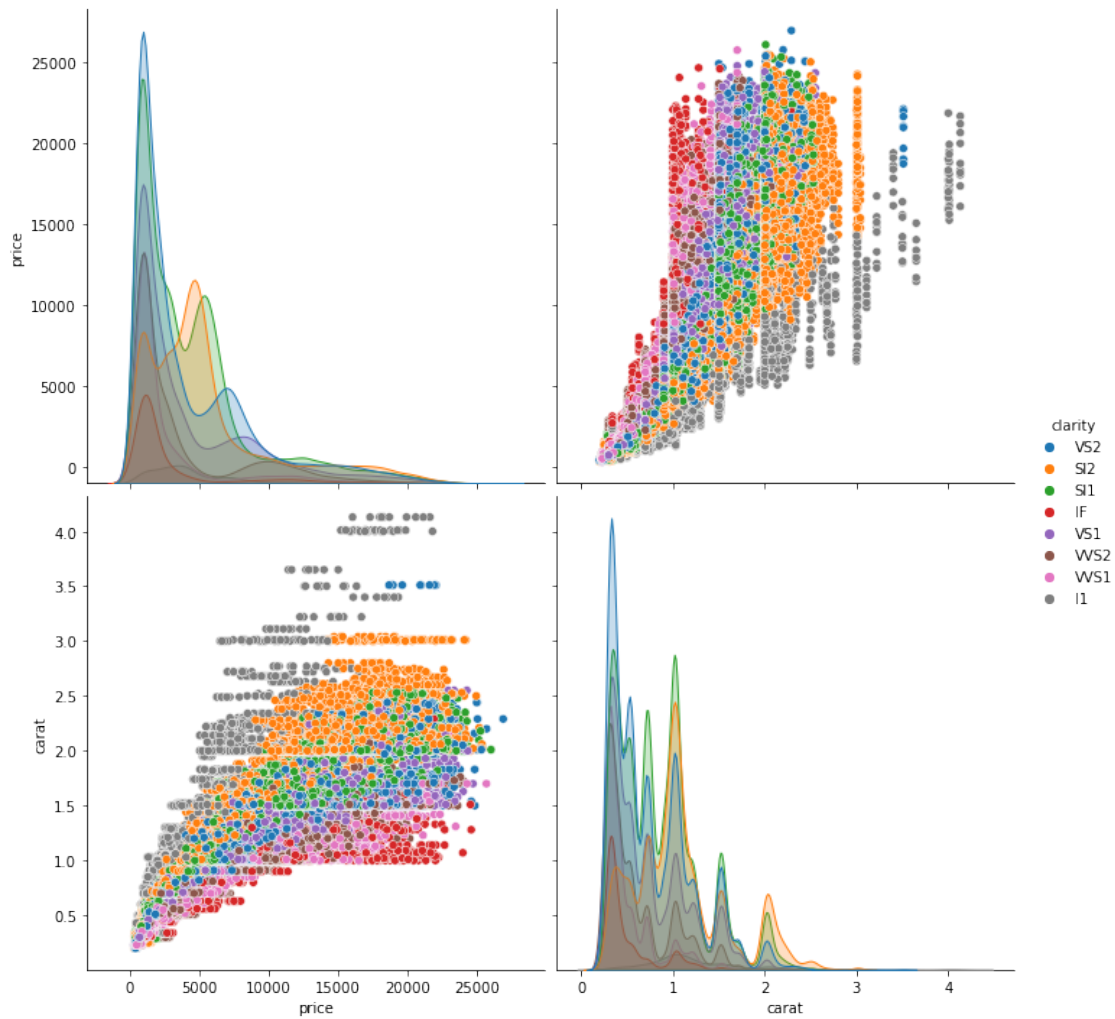
```
[37]: sns.pairplot(diamonds[["price", "carat", "color"]], hue = "color", height = 5)
plt.show()
sns.barplot(x = "carat", y = "color", data = diamonds)
plt.show()
sns.barplot(x = "price", y = "color", data = diamonds)
plt.show()
```







```
[38]: sns.pairplot(diamonds[["price", "carat", "clarity"]], hue = "clarity", height = 5)
```

```
plt.show()
sns.barplot(x = "carat", y = "clarity", data = diamonds)
plt.show()
sns.barplot(x = "price", y = "clarity", data = diamonds)
plt.show()
```





```
[39]: diamonds.describe()
```

```
[39]:
```

	carat	depth	table	price \
count	322553.000000	322553.000000	322553.000000	322553.000000
mean	0.797825	61.748797	57.458247	4396.330634
std	0.474806	1.434381	2.239518	4496.460572
min	0.200000	43.000000	43.000000	304.000000
25%	0.400000	61.000000	56.000000	1054.000000
50%	0.700000	61.800000	57.000000	2680.000000
75%	1.040000	62.500000	59.000000	5984.000000
max	4.130000	79.000000	95.000000	26930.000000

	x	y	z
count	322553.000000	322553.000000	322553.000000
mean	5.730354	5.732535	3.538676
std	1.122976	1.114397	0.713041
min	0.000000	0.000000	0.000000
25%	4.710000	4.720000	2.910000
50%	5.690000	5.710000	3.520000
75%	6.530000	6.530000	4.030000
max	10.140000	10.100000	31.800000

```
[40]: diamonds.head()
```

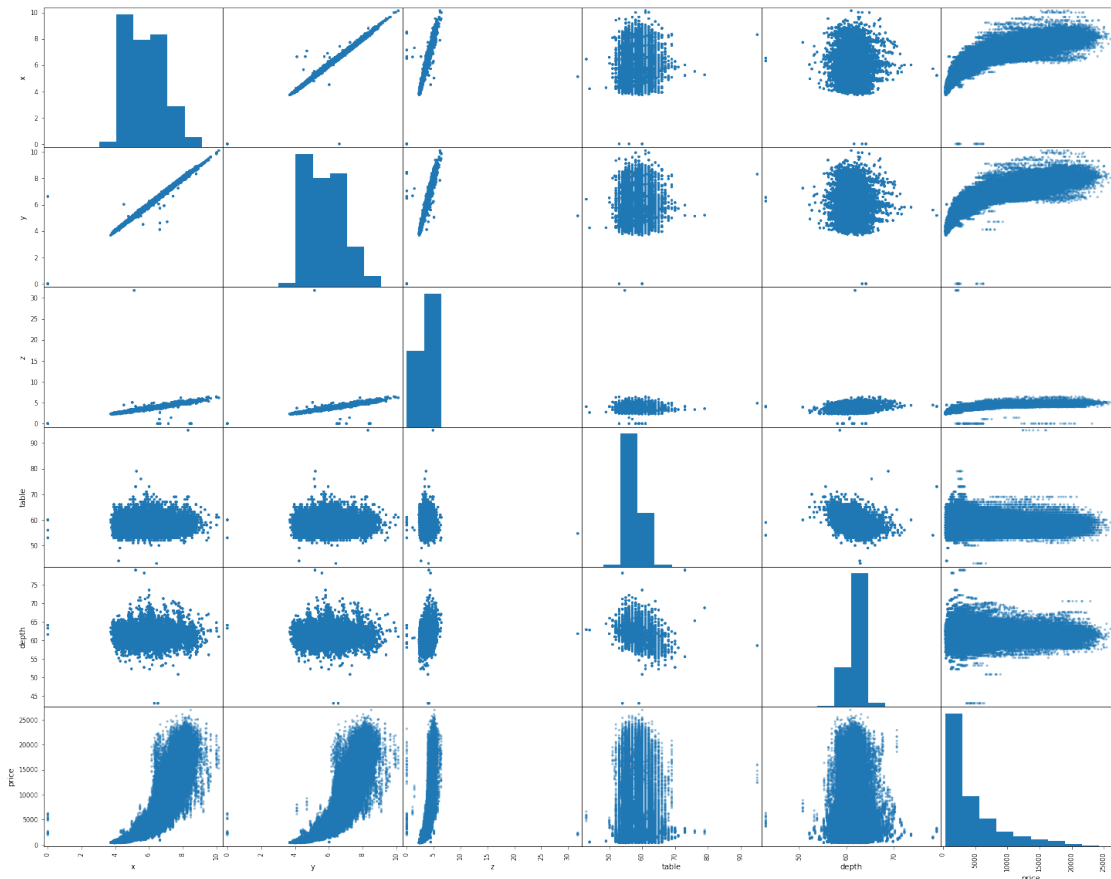
```
[40]:
```

	carat	cut	color	clarity	depth	table	price	x	y	z
79064	1.09	Ideal	F	VS2	61.2	56.0	7616.0	6.68	6.65	4.08
80834	1.50	Very Good	I	VS2	60.1	60.0	9339.0	7.38	7.35	4.43
106815	1.19	Premium	H	SI2	61.9	58.0	4525.0	6.84	6.77	4.21
269312	0.75	Premium	I	SI1	62.0	60.0	2744.0	5.80	5.75	3.57
8251	1.21	Premium	H	VS2	60.8	62.0	5461.0	6.82	6.78	4.13

```
[41]: diamonds.to_csv("C:/Users/Owner/Desktop/diamond_set/new_diamond.csv")
```

```
[42]: from pandas.plotting import scatter_matrix

attributes = ["x", "y", "z", "table", "depth", "price"]
scatter_matrix(diamonds[attributes], figsize=(25, 20))
plt.show()
```



```
[43]: # Do not stratify the label
diamonds = strat_train_set.drop("price", axis = 1)

# Set a new dataset label variable
diamond_labels = strat_train_set["price"].copy()

# Drop all the category, so we could have only numeric
diamonds_num = diamonds.drop(["cut", "color", "clarity"], axis = 1)
diamonds_num.head()
```

```
[43]:
```

	carat	depth	table	x	y	z
79064	1.09	61.2	56.0	6.68	6.65	4.08
80834	1.50	60.1	60.0	7.38	7.35	4.43
106815	1.19	61.9	58.0	6.84	6.77	4.21
269312	0.75	62.0	60.0	5.80	5.75	3.57
8251	1.21	60.8	62.0	6.82	6.78	4.13

```
[44]: from sklearn.preprocessing import StandardScaler
```

```
# Perform the feature scaling on the numeric attributes of the dataset
num_scaler = StandardScaler()
diamonds_num_scaled = num_scaler.fit_transform(diamonds_num)

# Preview
pd.DataFrame(diamonds_num_scaled).head()
```

```
[44]:
```

	0	1	2	3	4	5
0	0.615358	-0.382603	-0.651144	0.845652	0.823285	0.759178
1	1.478870	-1.149485	1.134957	1.468997	1.451428	1.250035
2	0.825971	0.105414	0.241906	0.988131	0.930967	0.941496
3	-0.100725	0.175130	1.134957	0.062019	0.015672	0.043931
4	0.868094	-0.661469	2.028008	0.970321	0.939940	0.829301

0.1.2 Build prediction model based on multiple features

```
[45]: from sklearn.preprocessing import StandardScaler

# Perform the feature scaling on the numeric attributes of the dataset
num_scaler = StandardScaler()
diamonds_num_scaled = num_scaler.fit_transform(diamonds_num)

# Preview
pd.DataFrame(diamonds_num_scaled).head()
```

```
[45]:
```

	0	1	2	3	4	5
0	0.615358	-0.382603	-0.651144	0.845652	0.823285	0.759178
1	1.478870	-1.149485	1.134957	1.468997	1.451428	1.250035
2	0.825971	0.105414	0.241906	0.988131	0.930967	0.941496
3	-0.100725	0.175130	1.134957	0.062019	0.015672	0.043931
4	0.868094	-0.661469	2.028008	0.970321	0.939940	0.829301

```
[46]: #Handling Catagorical Variables
# We need only the category attributes to work with here
diamonds_cat = diamonds[["cut", "color", "clarity"]]
diamonds_cat.head()
```

```
[46]:
```

	cut	color	clarity
79064	Ideal	F	VS2
80834	Very Good	I	VS2
106815	Premium	H	SI2
269312	Premium	I	SI1
8251	Premium	H	VS2

```
[47]: from sklearn.preprocessing import OneHotEncoder

# Perform the one-hot encoding on the category attributes of the dataset
```

```
cat_encoder = OneHotEncoder()
diamonds_cat_encoded = cat_encoder.fit_transform(diamonds_cat)

# Convert the encoded categories to arrays and Preview
pd.DataFrame(diamonds_cat_encoded.toarray()).head()
```

```
[47]:
```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	\
0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	
4	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	

	15	16	17	18	19
0	0.0	0.0	1.0	0.0	0.0
1	0.0	0.0	1.0	0.0	0.0
2	1.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	1.0	0.0	0.0

```
[48]: from sklearn.compose import ColumnTransformer

num_attribs = list(diamonds_num)
cat_attribs = ["cut", "color", "clarity"]

# Pipeline to transform our dataset
pipeline = ColumnTransformer([
    ("num", StandardScaler(), num_attribs), # Perform feaured scaling on
    ↪ numeric attributes
    ("cat", OneHotEncoder(), cat_attribs) # Perform One-Hot encoding on the
    ↪ category attributes
])
```

```
[49]: # Transformed dataset to feed the ML Algorithm
diamonds_ready = pipeline.fit_transform(diamonds)

# Preview
pd.DataFrame(diamonds_ready).head()
```

```
[49]:
```

	0	1	2	3	4	5	6	7	8	\
0	0.615358	-0.382603	-0.651144	0.845652	0.823285	0.759178	0.0	0.0	1.0	
1	1.478870	-1.149485	1.134957	1.468997	1.451428	1.250035	0.0	0.0	0.0	
2	0.825971	0.105414	0.241906	0.988131	0.930967	0.941496	0.0	0.0	0.0	
3	-0.100725	0.175130	1.134957	0.062019	0.015672	0.043931	0.0	0.0	0.0	
4	0.868094	-0.661469	2.028008	0.970321	0.939940	0.829301	0.0	0.0	0.0	

	9	...	16	17	18	19	20	21	22	23	24	25
--	---	-----	----	----	----	----	----	----	----	----	----	----

```

0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  1.0  0.0  0.0
1  0.0  ...  1.0  0.0  0.0  0.0  0.0  0.0  0.0  1.0  0.0  0.0
2  1.0  ...  0.0  0.0  0.0  0.0  0.0  1.0  0.0  0.0  0.0  0.0
3  1.0  ...  1.0  0.0  0.0  0.0  1.0  0.0  0.0  0.0  0.0  0.0
4  1.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  1.0  0.0  0.0

```

[5 rows x 26 columns]

```

[50]: from sklearn.metrics import mean_squared_error
      from sklearn.model_selection import cross_val_score
      from random import randint

      # Our test set
      # Remove target feature from test set
      X_test = strat_test_set.drop("price", axis = 1)
      # Have stand alone features
      y_test = strat_test_set["price"].copy()

      # Our models performance holder
      models_rmse = [] # Holds Models originalMSE
      cvs_rmse_mean = [] # Holds the Cross Validation RMSE Mean
      tests_rmse = [] # Holds the tests RMSE
      tests_accuracy = [] # Holds the tests accuracy
      models = [] # Holds the models name

      def display_model_performance(model_name, model, diamonds = diamonds_ready,
      ↪ labels = diamond_labels,
      ↪ models_rmse = models_rmse, cvs_rmse_mean =
      ↪ cvs_rmse_mean, tests_rmse = tests_rmse,
      ↪ tests_accuracy = tests_accuracy, pipeline =
      ↪ pipeline, X_test = X_test,
      ↪ y_test = y_test, cv = True):
          # Fit dataset in model
          model.fit(diamonds, labels)

          # Setup predictions
          predictions = model.predict(diamonds)

          # Get models performance
          model_mse = mean_squared_error(labels, predictions)
          model_rmse = np.sqrt(model_mse)

          # Cross validation
          cv_score = cross_val_score(model, diamonds, labels, scoring =
          ↪ "neg_mean_squared_error", cv = 10)

```

```

cv_rmse = np.sqrt(-cv_score)
cv_rmse_mean = cv_rmse.mean()

print("RMSE: %.4f" %model_rmse)
models_rmse.append(model_rmse)

print("CV-RMSE: %.4f" %cv_rmse_mean)
cvs_rmse_mean.append(cv_rmse_mean)

print("--- Test Performance ---")

X_test_prepared = pipeline.transform(X_test)

# Fit test dataset in model
model.fit(X_test_prepared, y_test)

# Setup test predictions
test_predictions = model.predict(X_test_prepared)

# Get models performance on test
test_model_mse = mean_squared_error(y_test, test_predictions)
test_model_rmse = np.sqrt(test_model_mse)
print("RMSE: %.4f" %test_model_rmse)
tests_rmse.append(test_model_rmse)

# Tests accuracy
test_accuracy = round(model.score(X_test_prepared, y_test) * 100, 2)
print("Accuracy:", str(test_accuracy)+"%")
tests_accuracy.append(test_accuracy)

# Check how well model works on Test set by comparing prices
start = randint(1, len(y_test))
some_data = X_test.iloc[start:start + 7]
some_labels = y_test.iloc[start:start + 7]
some_data_prepared = pipeline.transform(some_data)
print("Predictions:\t", model.predict(some_data_prepared))
print("Labels:\t\t", list(some_labels))

models.append(model_name)

# Preview plot
plt.scatter(diamond_labels, model.predict(diamonds_ready))
plt.xlabel("Actual")
plt.ylabel("Predicted")
x_lim = plt.xlim()
y_lim = plt.ylim()

```

```
plt.plot(x_lim, y_lim, "k--")
plt.show()

print("----- Test -----")
plt.scatter(y_test, model.predict(X_test_prepared))
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.plot(x_lim, y_lim, "k--")
plt.show()
```


We can now start fitting models and get their performance error. We are using Root Mean Squared Error for our performance measure.

0.2 Linear Regression


```
[51]: from sklearn.linear_model import LinearRegression

lin_reg = LinearRegression(normalize = True)
display_model_performance("Linear Regression", lin_reg)
```

```
RMSE: 1372.0532
CV-RMSE: 1372.2307
--- Test Performance ---
RMSE: 1376.8318
Accuracy: 90.73%
Predictions:      [ 1036.3606775   2871.16259804   122.81520923 13419.49156411
  1017.92549942  3454.78791488  1347.93573761]
Labels:          [644.0, 3110.0, 722.0, 14433.0, 831.0, 2906.0, 843.0]
```

----- Test -----



0.3 Decision Tree Regression

```
[52]: from sklearn.tree import DecisionTreeRegressor
```

```
tree_reg = DecisionTreeRegressor(random_state = 42)  
display_model_performance("Decision Tree Regression", tree_reg)
```

RMSE: 510.7221

CV-RMSE: 573.9419


--- Test Performance ---

RMSE: 425.3440


Accuracy: 99.12%

Predictions: [441. 2532.25 8786.6 3430.33333333 8069.
2515. 6425.5]

Labels: [441.0, 2557.0, 8296.0, 3304.0, 8069.0, 2491.0, 6293.0]



----- Test -----



0.3.1 Random Forest Regression

```
[53]: from sklearn.ensemble import RandomForestRegressor
```

```
forest_reg = RandomForestRegressor(n_estimators = 10, random_state = 42)
display_model_performance("Random Forest Regression", forest_reg)
```

RMSE: 513.7440

CV-RMSE: 577.4522


--- Test Performance ---

RMSE: 455.3055


Accuracy: 98.99%

Predictions: [3790.48333333 919.65166667 5636.62357143 20116.74833333
1000.49166667 6283.13333333 4879.45]

Labels: [3909.0, 1021.0, 5792.0, 20072.0, 1007.0, 6452.0, 4482.0]



----- Test -----



0.4 Ridge Regression

```
[54]: from sklearn.linear_model import Ridge
```

```
ridge_reg = Ridge(normalize = True)  
display_model_performance("Ridge Regression", ridge_reg)
```

RMSE: 2043.7825

CV-RMSE: 2043.9321


--- Test Performance ---

RMSE: 2059.8497


Accuracy: 79.26%

Predictions: [1240.11985888 4294.6933905 6012.59313031 11976.15806539
11343.10027794 526.42254813 2783.02404763]

Labels: [1025.0, 2850.0, 4422.0, 21233.0, 17192.0, 680.0, 1412.0]



----- Test -----



0.5 Lasso Regression

```
[55]: from sklearn.linear_model import Lasso
```

```
lasso_reg = Lasso(normalize = True)
display_model_performance("Lasso Regression", lasso_reg)
```

RMSE: 1898.2986

CV-RMSE: 1889.7609

--- Test Performance ---


RMSE: 1732.4304

Accuracy: 85.33%


Predictions: [815.15575328 2310.74823845 6080.5518979 981.33269608
2476.92518125

6548.26027978 3338.63234175]

Labels: [558.0, 2043.0, 4177.0, 996.0, 1835.0, 11860.0, 2959.0]



----- Test -----



0.6 Elastic Net Regression

```
[56]: from sklearn.linear_model import ElasticNet
```

```
net_reg = ElasticNet()  
display_model_performance("Elastic Net Regression", net_reg)
```

RMSE: 1987.6651

CV-RMSE: 1987.7956


--- Test Performance ---

RMSE: 2004.1156


Accuracy: 80.36%

Predictions: [5668.43803649 1050.13153585 3117.56088973 239.48574925
6472.05781275 11379.39931957 6828.97197001]

Labels: [4851.0, 1250.0, 2055.0, 837.0, 7021.0, 18377.0, 6147.0]



----- Test -----



0.7 Save the model

```
[57]: #create Pickle file
import pickle
pickle_out = open("forest_reg.pkl","wb")
pickle.dump(forest_reg,pickle_out)
pickle_out.close()
```

0.8 Model Deployment of the prediction using Streamlit

```
[58]: #import streamlit as st
```

```
[59]: #loading the trained model
#pickle_in = open('model.pkl', 'rb')
#regressor = pickle.load(pickle_in)
```

```
[60]: #diamonds.to_csv("model_prediction_dataset.csv")
```

```
[61]: #diamonds.to_csv("C:/Users/Owner/Downloads/prediction_model.csv")
```

```
[ ]:
```